

RHEL STIG RAG - Podman Deployment Guide

Comprehensive guide for deploying the RHEL STIG RAG system using Podman with container orchestration and systemd integration.

Overview

This Podman deployment provides:

- **Rootless Security:** Enhanced security through rootless containers
- **Systemd Integration:** Native systemd service management
- **Pod Orchestration:** Multi-container pod support
- **Auto-recovery:** Automatic restart and health monitoring
- **RHEL-Native:** Optimized for Red Hat environments

Quick Start

Option 1: Automated Script

```
bash

# Clone repository
git clone <repository-url>
cd rhel-stig-rag

# Make script executable
chmod +x deploy-podman.sh

# Deploy (rootless mode - default)
./deploy-podman.sh deploy

# Deploy in rootful mode
./deploy-podman.sh --rootful deploy
```

Option 2: Ansible Deployment

```
bash
```

```
# Deploy with Ansible
```

```
cd ansible
```

```
ansible-playbook -i inventory/hosts podman-deploy.yml
```

```
# Deploy in rootful mode
```

```
ansible-playbook -i inventory/hosts podman-deploy.yml -e "stig_rag_podman_mode=rootful"
```

Prerequisites

System Requirements

- **OS:** RHEL 8+, CentOS 8+, Fedora 35+, Ubuntu 20.04+
- **Memory:** 4GB RAM minimum (8GB recommended)
- **Storage:** 5GB free space
- **CPU:** 2 cores minimum

Software Requirements

```
bash
```

```
# RHEL/CentOS/Fedora
```

```
sudo dnf install podman buildah skopeo
```

```
# Ubuntu/Debian
```

```
sudo apt install podman buildah
```

Deployment Modes

Rootless Mode (Recommended)

Advantages:

- Enhanced security (no root privileges)
- User namespace isolation
- No daemon running as root

Usage:

```
bash
```

```
# Deploy rootless
```

```
./deploy-podman.sh --rootless deploy
```

```
# Status check
```

```
podman pod ps
```

```
systemctl --user status stig-rag-pod
```

Rootful Mode

Advantages:

- Better performance for I/O operations
- Easier networking configuration
- System-wide service management

Usage:

```
bash
```

```
# Deploy rootful
```

```
./deploy-podman.sh --rootful deploy
```

```
# Status check
```

```
sudo podman pod ps
```

```
sudo systemctl status stig-rag-pod
```

Container Architecture

Pod Structure

```
stig-rag-pod
```

```
├─ stig-rag (main application)
```

```
├─ stig-rag-metrics (optional monitoring)
```

```
└─ shared network namespace
```

Volume Mapping

Host Directory	→ Container Path
~/stig-rag-data/stig_data	→ /app/stig_data
~/stig-rag-data/stig_chroma_db	→ /app/stig_chroma_db
~/stig-rag-logs	→ /app/logs
~/stig-rag-config	→ /app/config

Configuration

Environment Variables

Key configuration options:

```
bash

# Application settings
APP_HOST=0.0.0.0
APP_PORT=8000
DEFAULT_RHEL_VERSION=9

# Performance tuning
CHUNK_SIZE=1000
EMBEDDING_MODEL=all-MiniLM-L6-v2
LLM_PROVIDER=huggingface

# Resource limits
MEMORY_LIMIT=4g
CPU_LIMIT=2.0
```

Custom Configuration

Edit the configuration file:

```
bash

# Rootless mode
vim ~/stig-rag-config/config.env

# Rootful mode
sudo vim /etc/stig-rag/config.env

# Restart to apply changes
./manage-stig-rag.sh restart
```

Management Commands

Using Management Script

```
bash

# Check status
./manage-stig-rag.sh status

# View Logs
./manage-stig-rag.sh logs
./manage-stig-rag.sh logs -f # Follow Logs

# Start/stop services
./manage-stig-rag.sh start
./manage-stig-rag.sh stop
./manage-stig-rag.sh restart

# Health check
./manage-stig-rag.sh health

# Update containers
./manage-stig-rag.sh update

# Open shell in container
./manage-stig-rag.sh shell

# Manual backup
./manage-stig-rag.sh backup
```

Using Podman Commands

```
bash
```

```
# Pod management
```

```
podman pod ps
```

```
podman pod start stig-rag-pod
```

```
podman pod stop stig-rag-pod
```

```
# Container management
```

```
podman ps
```

```
podman logs stig-rag
```

```
podman exec -it stig-rag /bin/bash
```

```
# Health check
```

```
podman healthcheck run stig-rag
```

Using Systemd Commands

Rootless mode:

```
bash
```

```
systemctl --user status stig-rag-pod
```

```
systemctl --user restart stig-rag-pod
```

```
systemctl --user enable stig-rag-pod
```

Rootful mode:

```
bash
```

```
sudo systemctl status stig-rag-pod
```

```
sudo systemctl restart stig-rag-pod
```

```
sudo systemctl enable stig-rag-pod
```

Advanced Features

Systemd Integration with Quadlet

For systemd 250+ (RHEL 9+), you can use Quadlet for better integration:

```
bash
```

```
# Enable Quadlet mode
```

```
ansible-playbook podman-deploy.yml -e "stig_rag_podman_quadlet_enabled=true"
```

```
# Check generated services
```

```
systemctl --user list-units | grep stig-rag
```

Auto-Updates

Enable automatic container updates:

```
bash
```

```
# Enable auto-update
```

```
ansible-playbook podman-deploy.yml -e "stig_rag_podman_auto_update=true"
```

```
# Check auto-update status
```

```
systemctl --user status podman-auto-update.timer
```

Monitoring Integration

Deploy with Prometheus metrics:

```
bash
```

```
# Deploy with monitoring
```

```
ansible-playbook podman-deploy.yml -e "stig_rag_podman_expose_metrics=true"
```

```
# Access metrics
```

```
curl http://localhost:9090/metrics
```

Backup Management

Automated backups are configured by default:

```
bash
```

```
# Check backup timer
```

```
systemctl --user status stig-rag-backup.timer
```

```
# Manual backup
```

```
systemctl --user start stig-rag-backup.service
```

```
# Restore from backup
```

```
tar -xzf ~/stig-rag-backups/stig-rag-backup-*.tar.gz -C ~/stig-rag-data/
```

Networking

Port Configuration

Default ports:

- **8000**: STIG RAG API
- **9090**: Prometheus metrics (if enabled)

Custom Network

The deployment creates a dedicated network:

```
bash
```

```
# Inspect network
```

```
podman network inspect stig-rag-net
```

```
# Connect additional containers
```

```
podman run --network stig-rag-net <other-container>
```

Security Considerations

SELinux Integration

Podman automatically handles SELinux contexts with the `:Z` flag:

```
bash
```

```
# Check SELinux contexts
```

```
ls -Z ~/stig-rag-data/
```

Security Options

Default security configurations:

- Dropped capabilities: ALL
- Security labels: container_runtime_t
- Read-only root filesystem (optional)
- PID limits

Rootless Benefits

- No root daemon
- User namespace isolation
- Limited filesystem access
- Reduced attack surface

Troubleshooting

Common Issues

1. Permission Denied

Rootless mode issue:

```
bash

# Check user namespace
podman unshare cat /proc/self/uid_map

# Fix storage permissions
podman system reset
```

2. Port Already in Use

```
bash

# Change port
./deploy-podman.sh --port 8080 deploy

# Or stop conflicting service
sudo ss -tulpn | grep :8080
```

3. Container Won't Start

```
bash
```

```
# Check Logs
```

```
podman logs stig-rag
```

```
# Check systemd
```

```
systemctl --user status stig-rag-pod -l
```

```
journalctl --user -u stig-rag-pod -f
```

4. Health Check Failing

```
bash
```

```
# Manual health check
```

```
curl -v http://localhost:8000/health
```

```
# Check container health
```

```
podman inspect stig-rag | jq '[0].State.Health'
```

Debug Mode

Enable debug logging:

```
bash
```

```
# Deploy with debug
```

```
ansible-playbook podman-deploy.yml -e "stig_rag_log_level=DEBUG"
```

```
# Check debug logs
```

```
./manage-stig-rag.sh logs | grep DEBUG
```

Performance Tuning

Resource Limits

Adjust based on your system:

```
yaml
```

```
# In vars/podman.yml
```

```
stig_rag_podman_memory_limit: "8g"
```

```
stig_rag_podman_cpu_limit: "4.0"
```

```
stig_rag_podman_pids_limit: 2048
```

Storage Optimization

```
bash

# Monitor storage usage
podman system df

# Clean up unused resources
podman system prune -a

# Optimize storage driver
# Edit ~/.config/containers/storage.conf
```

Model Optimization

Use smaller models for better performance:

```
yaml

stig_rag_embedding_model: "all-MiniLM-L6-v2" # Faster, smaller
stig_rag_llm_model: "microsoft/DialoGPT-small" # Reduced memory
```

Integration Examples

CI/CD Pipeline

```
yaml

# .gitlab-ci.yml
deploy:
  script:
    - ansible-playbook ansible/podman-deploy.yml
  only:
    - main
```

Load Balancer

```
bash
```

```
# Deploy multiple instances
```

```
for i in {1..3}; do
```

```
    ./deploy-podman.sh --port $((8000 + i)) --pod-name stig-rag-pod-$i deploy
```

```
done
```

```
# Configure nginx load balancer
```

```
# (nginx configuration not included)
```

Development Workflow

```
bash
```

```
# Development deployment
```

```
./deploy-podman.sh --rootless deploy
```

```
# Code changes
```

```
vim rhel_stig_rag.py
```

```
# Rebuild and update
```

```
./deploy-podman.sh build
```

```
./manage-stig-rag.sh update
```

Migration from Docker

Docker to Podman

```
bash
```

```
# Convert docker-compose.yml to Podman
```

```
podman-compose up -d
```

```
# Or use the provided Podman deployment
```

```
./deploy-podman.sh deploy
```

Volume Migration

```
bash
```

```
# Copy data from Docker volumes
```

```
docker cp stig-rag-container:/app/stig_data ~/stig-rag-data/
```

```
docker cp stig-rag-container:/app/stig_chroma_db ~/stig-rag-data/
```

Maintenance

Regular Updates

```
bash

# Update system packages
sudo dnf update podman

# Update container images
./manage-stig-rag.sh update

# Update STIG data
# (automation via ansible playbook)
ansible-playbook update_stig.yml
```

Monitoring Health

```
bash

# Check service health
./manage-stig-rag.sh health

# Monitor Logs
./manage-stig-rag.sh logs -f

# System resource usage
podman stats stig-rag
```

Backup and Recovery

```
bash

# Automated backup (daily)
systemctl --user enable stig-rag-backup.timer

# Manual backup
./manage-stig-rag.sh backup

# Recovery process
./deploy-podman.sh stop
tar -xzf backup.tar.gz -C ~/stig-rag-data/
./deploy-podman.sh start
```

Support and Resources

Useful Commands

```
bash
```

```
# Complete system info
```

```
podman info
```

```
# Network information
```

```
podman network ls
```

```
# Volume information
```

```
podman volume ls
```

```
# System events
```

```
podman events
```

Log Locations

- **Container logs:** `podman logs stig-rag`
- **Systemd logs:** `journalctl --user -u stig-rag-pod`
- **Application logs:** `~/stig-rag-logs/stig_rag.log`

Performance Monitoring

```
bash
```

```
# Resource usage
```

```
podman stats stig-rag
```

```
# Container inspection
```

```
podman inspect stig-rag
```

```
# Health check history
```

```
podman healthcheck run stig-rag
```

This Podman deployment provides enterprise-grade container management with enhanced security, better integration with RHEL/systemd, and comprehensive automation through Ansible. The rootless mode is particularly valuable for security-conscious environments while maintaining full functionality.